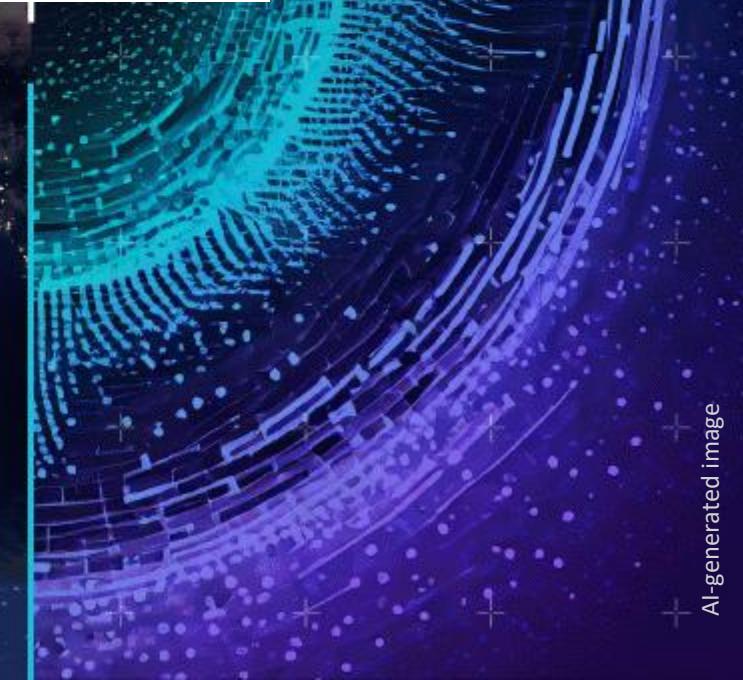
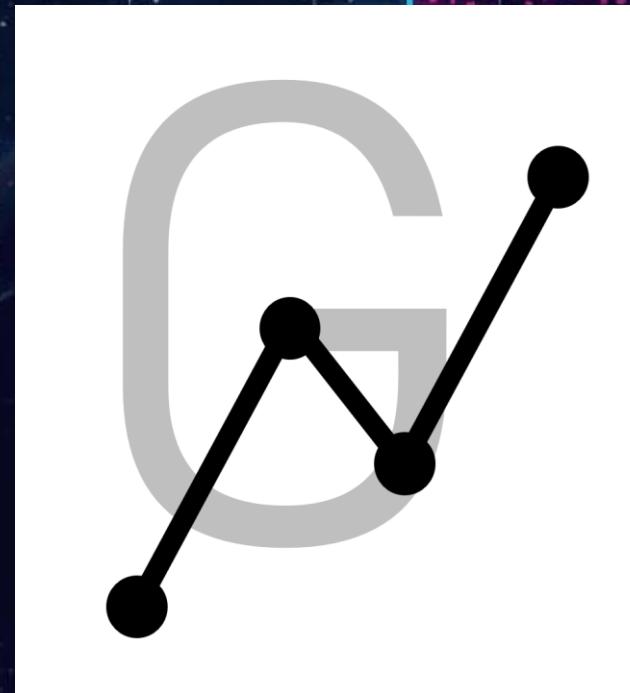


# arm

## NanoGraph

A graph interpreter for  
DSP/ML stream-based processing



# Summary : DSP/ML is complex, and this slows time-to-market

Complexity from the real-world “physical domains” and the DSP/ML “software computation domains”

- ⇒ Split the problem in smaller pieces (“computing nodes”)
- ⇒ nodes are stored in Flash and activated by an interpreted stream-processing scenario
- ⇒ A low-code approach enables easy addition of new nodes



April 2024

Your Global IoT Market Research Partner

## Time to market for IoT-connected products



Note: \*Time to market = time needed (in months) to get from project kick-off to first paying customer.

Source: IoT Analytics Research 2024 – IoT Commercialization & Business Model Adoption Report 2024. We welcome republishing of images but ask for source citation with a link to the original post and company website.



# NanoGraph vs. Traditional Firmware/Monolithic DSP/ML

Aspect	Graph Interpreter Approach	Traditional Firmware / Monolithic DSP/ML
<b>Modularity</b>	Highly modular; nodes are independent, reusable, and updatable	Monolithic; tightly coupled code, hard to reuse or update
<b>Extensibility</b>	Add new nodes/functions with minimal code (“low-code”)	Adding features requires significant code rewrite
<b>Portability</b>	Portable across hardware platforms; standard interfaces	Hardware-specific; porting requires major adaptation
<b>Time-to-Market</b>	Faster: rapid prototyping and deployment with pre-built nodes	Slower: longer development and testing cycles
<b>Maintainability</b>	Update or patch individual nodes without touching the whole system	Updates may require full firmware rebuild
<b>Security &amp; Reliability</b>	Interpreter-based execution reduces update risks; supports self-recovery	Updates can fail, potentially bricking devices
<b>Resource Efficiency</b>	Optimized memory management; small footprint	Often less efficient; higher memory/CPU usage
<b>ML/DSP Integration</b>	Native support for DSP/ML nodes, sensor fusion, and AI at the edge	Integration is manual and complex
<b>Customization</b>	End-users/integrators can customize graphs without deep firmware knowledge	Customization requires firmware expertise



# Splitting the problem, by looking at the users' focus

## Silicon vendors

Primary interest: Demonstrating the architecture's capabilities on micro-kernels

Other interests : the software ecosystem can switch to a new HW architecture

## Software developers

First interest : ease of development with tools, tutorials, compute libraries

Other interests : portability and performance scaling of existing code base

## System integrators

First interest : wide catalog of applications, multi-source, decent performance

Other interests : accelerate time to market with good development tools

## Key selling points of NanoGraph

Many "nodes" pre-installed  
in the Flash  
Standard interfaces, secured "Stores"  
libs : NEON/MVE.. malloc for TCM  
Language independent

Graph portability, AI done locally  
Low-code, Fast tuning  
Lower the risk of failed FW updates  
Self-recovery (drift, warm-boot)

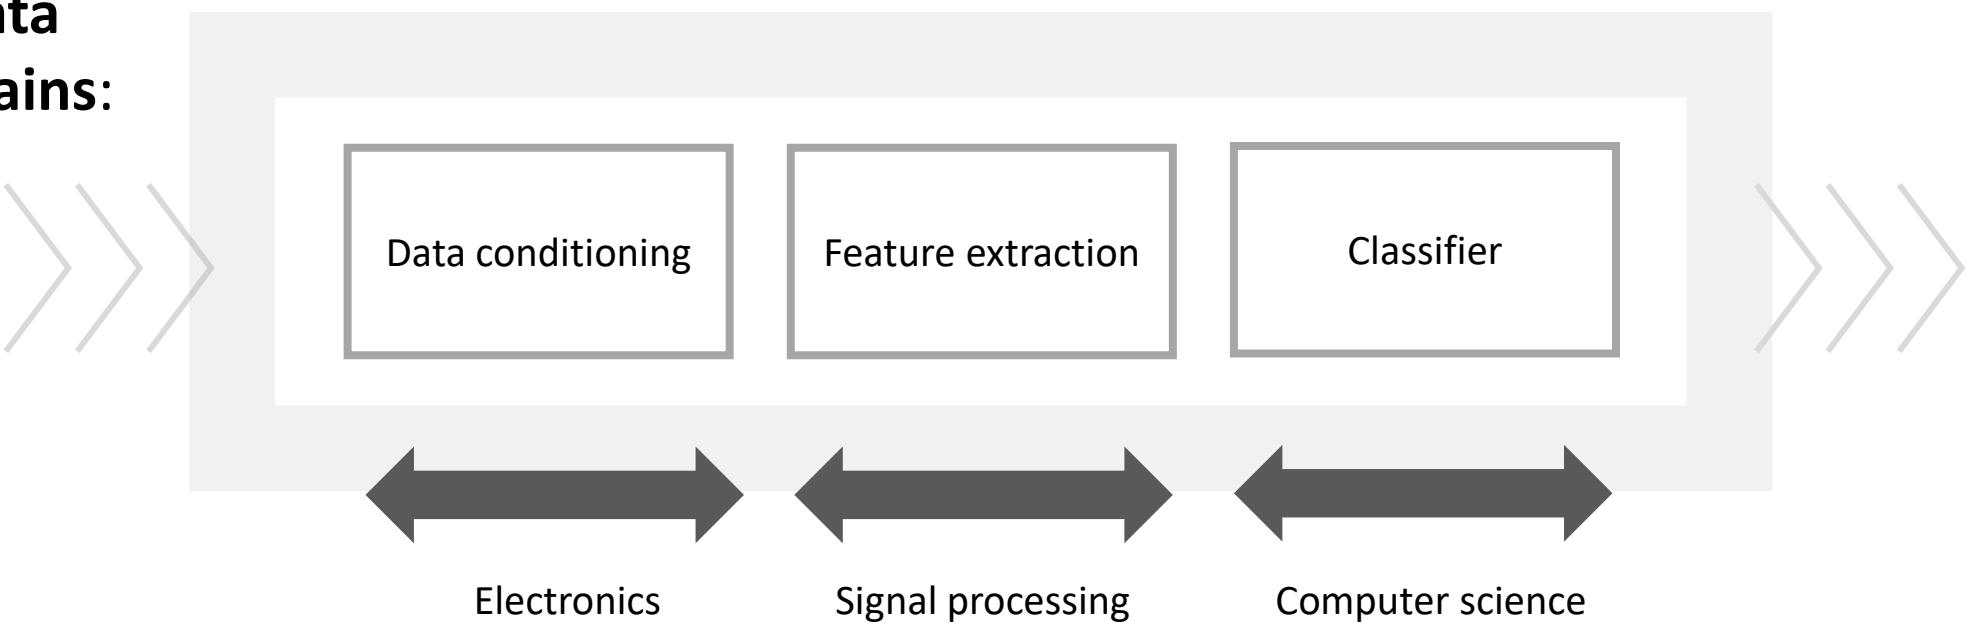


# Stream-based processing - different domains of expertise

The complexity of DSP/ML processing graphs stems from the need to integrate expertise across multiple domains

## Different **data physical domains**:

- Acoustics
- Electrical
- Chemical
- Mechanical
- ...

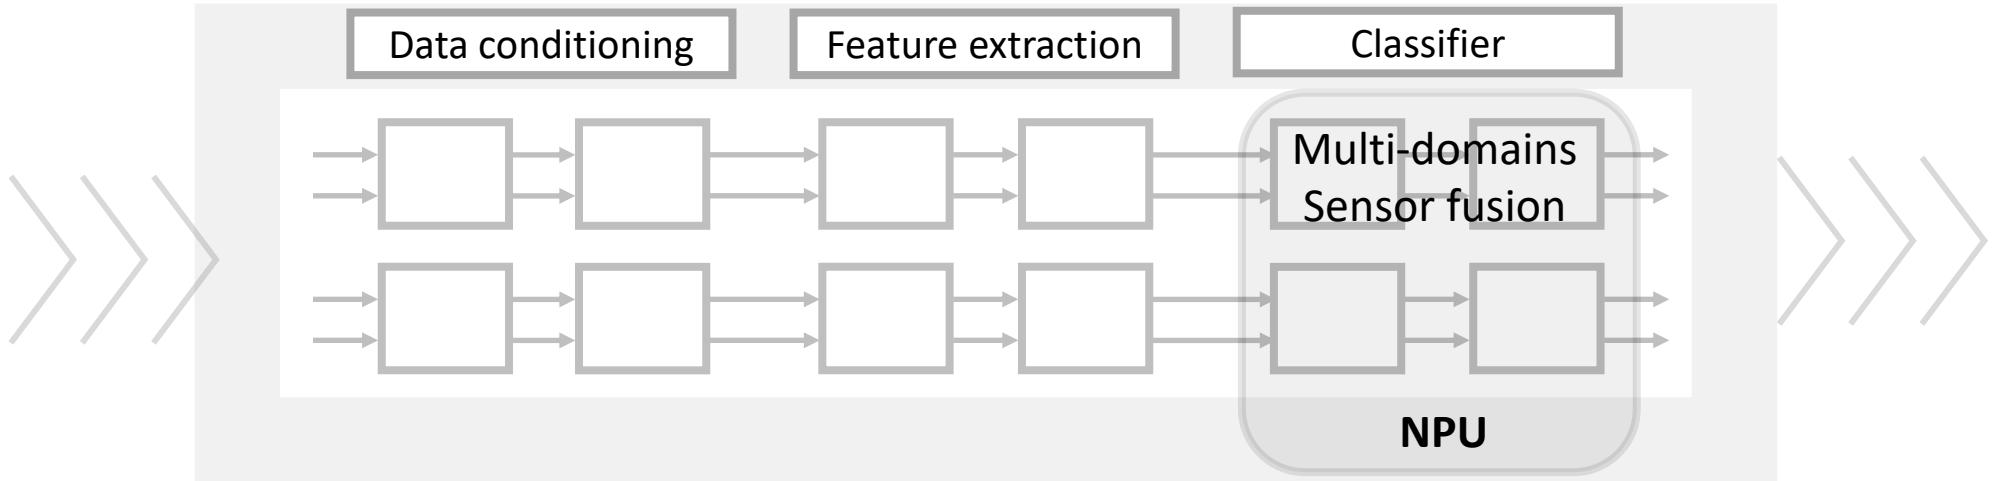


## Different **software engineering domains**



# Stream-based processing with graph of computing nodes

- Acoustics
- Electrical
- Chemical
- Mechanical
- ...

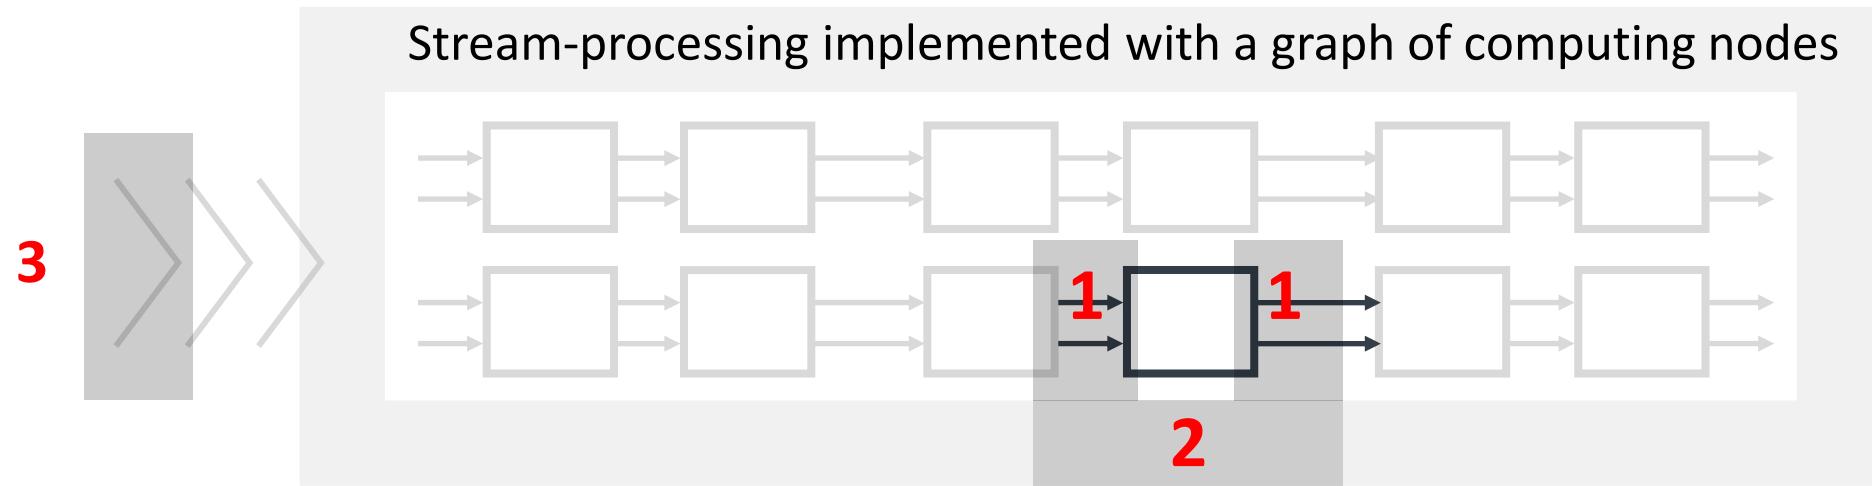


Our proposal : stream-processing is implemented with a graph of computing nodes **designed independently** (from different providers), some nodes can be pre-installed in the Flash of the platform manufacturer

The proof of concept is in production with the graph of [EEMBC audiomark](#) using four DSP nodes (beamformer, echo and noise suppressor and a [classifier node](#) for Key Word Spotting) running with or without NPU, but with the same node's interface

# Manifests of interfaces for nodes, Graph-I/O, Processor

Interfaces between nodes, the scheduler, and graph I/Os are standardized and formalized

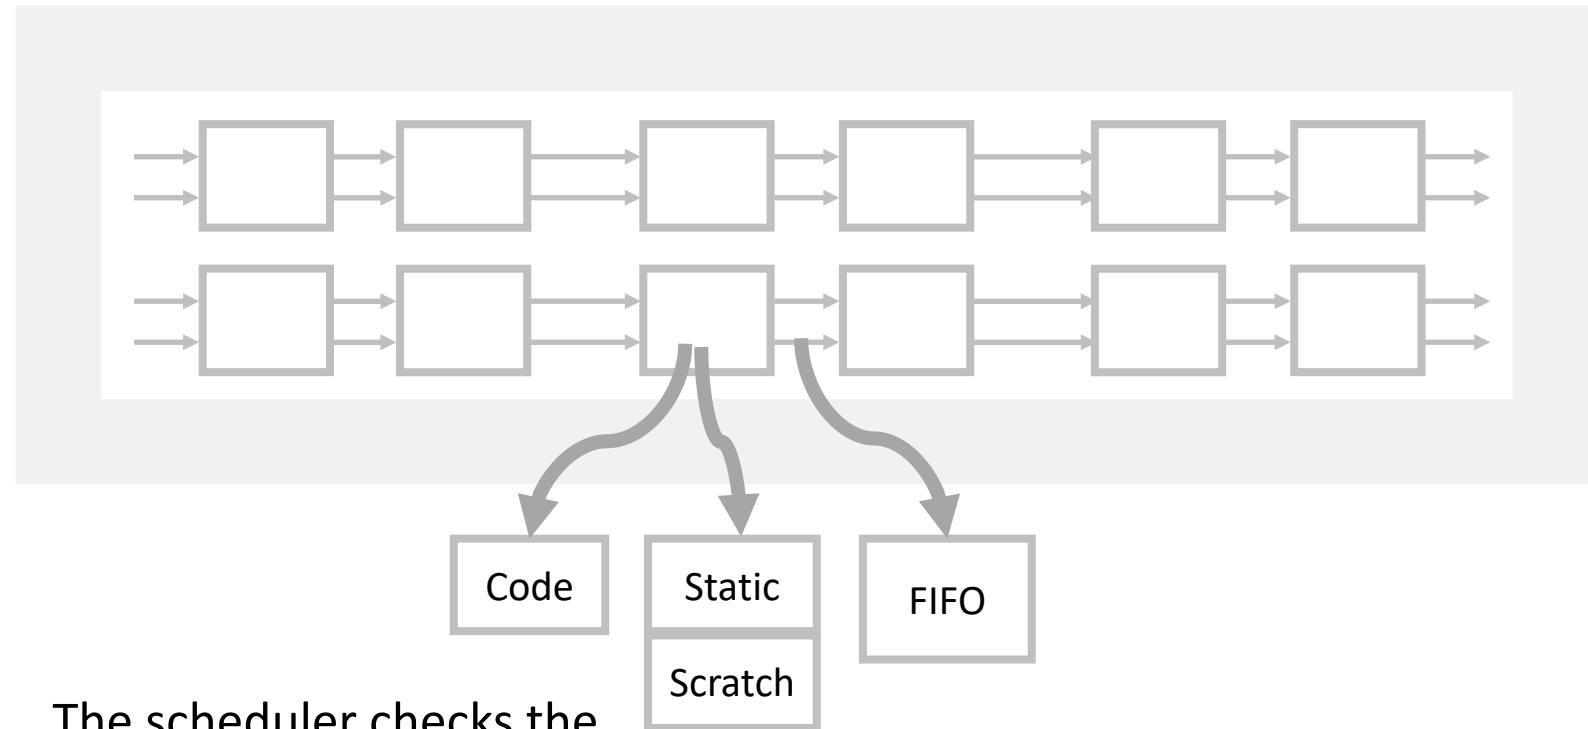


- 1 Inter-node interface** : data format (sampling rate, interleaving, raw format, frame size)
- 2 Processor interface** with nodes : memory allocation and TCM, compute libraries and NPU
- 3 Graph-I/O interfaces** : buffering and polling scheme, mixed-signal configuration of the domains

# NanoGraph interpreter and scheduler

The graph intermediate format is a text file, “compiled” to build a scheduling table and a memory mapping

The compiled graph is  
a linked list  
referencing memory  
buffers and node  
addresses

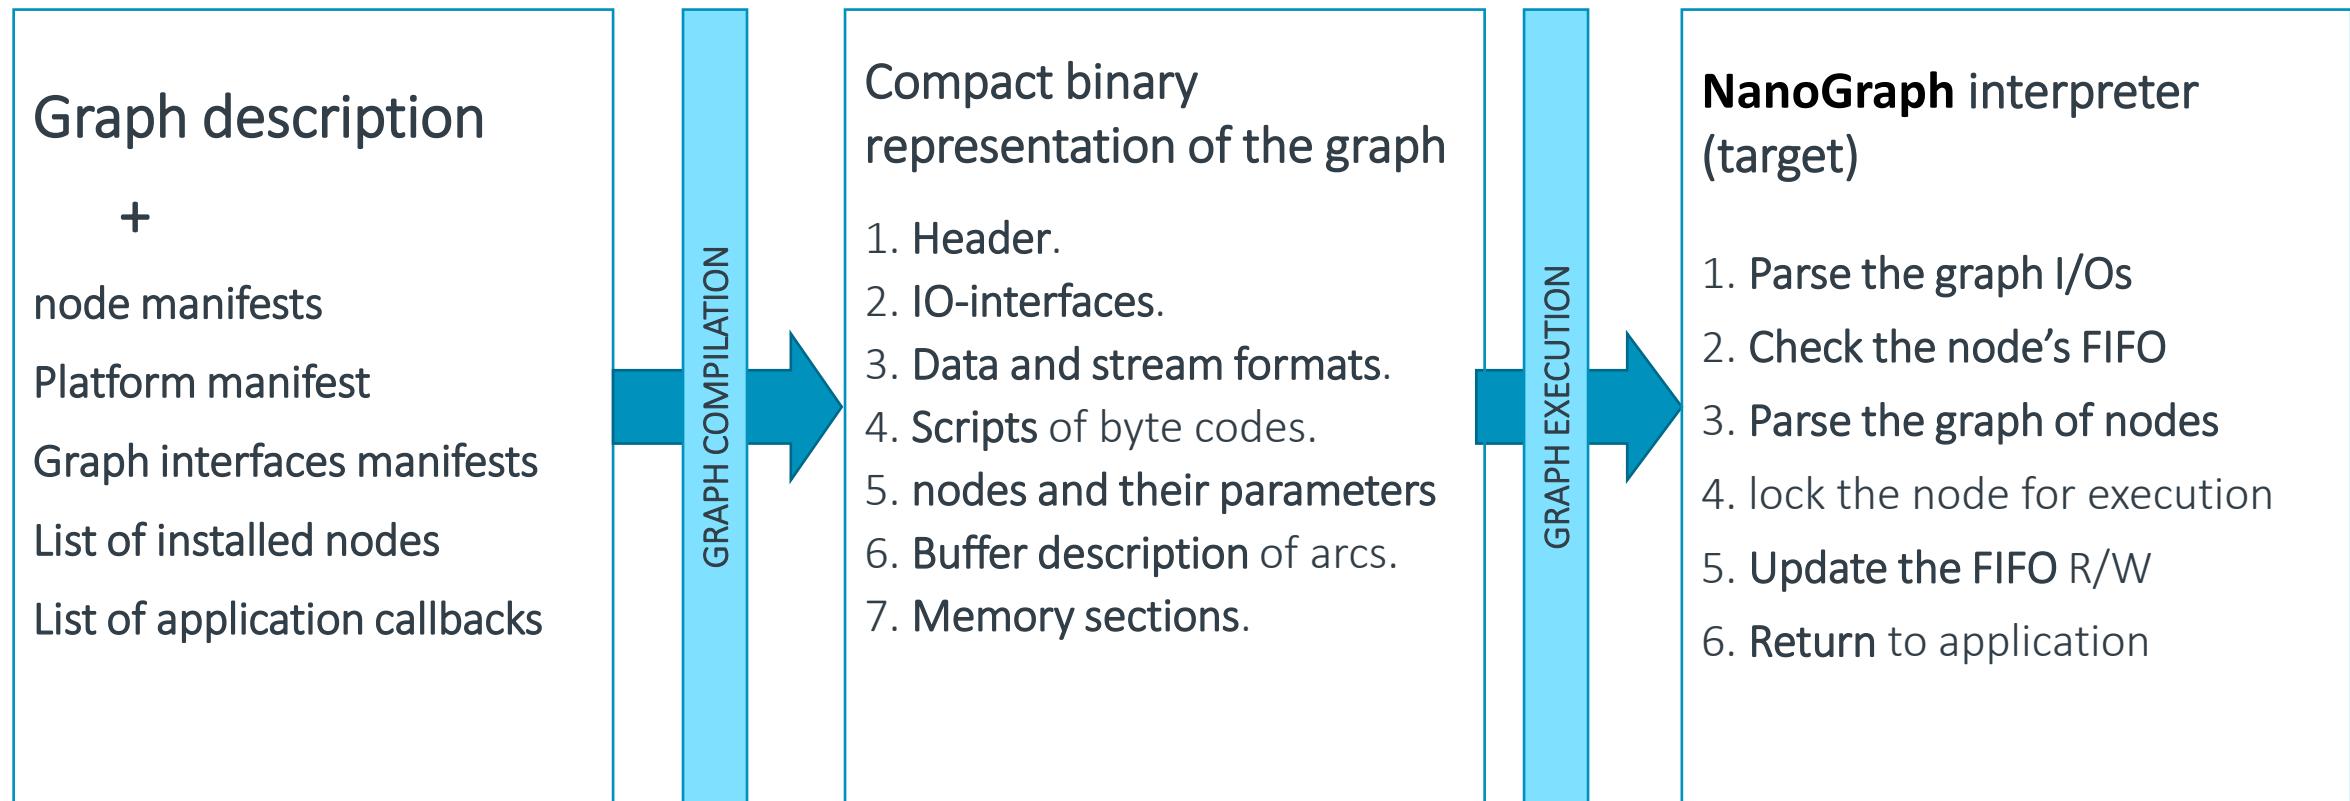


The scheduler checks the  
FIFO buffers before  
calling a node instance



# Compilation process using “Manifests”

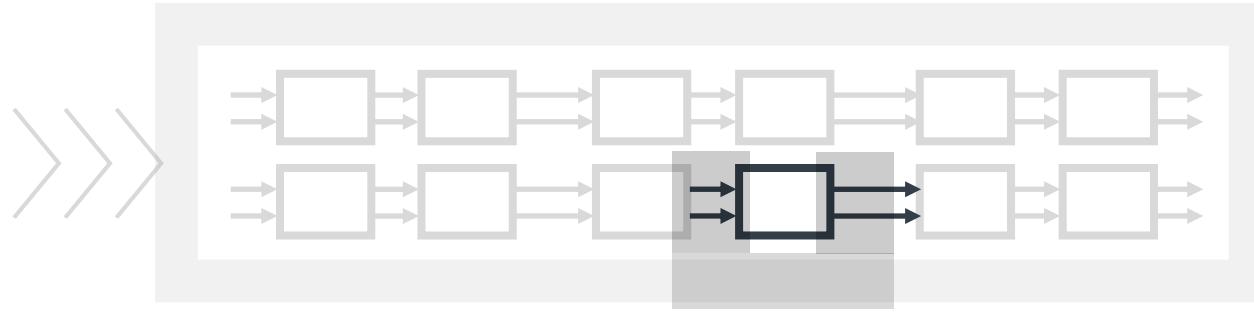
I/O and node manifests help the graph compiler build the memory map and define data flow between nodes.



# arm

## Graph design

# Manifests for nodes



## 1 Inter-node interface and interface with the platform :

a text file (readable syntax)  
done once at node delivery

```
; -----  
; SOFTWARE COMPONENT MANIFEST - "arm_stream_filter"  
;  
;  
node_developer_name    ARM           ; developer name  
node_name               arm_stream_filter ; node name  
  
node_using_arc_format  1             ; to let filter manage q15 and fp32  
node_mask_library       64            ; dependency with DSP services  
  
;  
;   MEMORY ALLOCATIONS  
  
node_mem                0             ; first memory bank (node instance)  
node_mem_alloc          76            ; amount of bytes  
  
node_mem                1             ; second memory bank (node fast working area)  
node_mem_alloc          52            ;  
node_mem_type           1             ; working memory  
node_mem_speed          2             ; critical fast  
;  
;   ARCS CONFIGURATION  
node_arc                0             ;  
node_arc_nb_channels     {1 1 2}      ; arc interleaved, options for the number of channels  
node_arc_raw_format      {1 2 1}      ; options for the raw format STREAM_S16, STREAM_FP32  
  
node_arc                1             ;  
node_arc_nb_channels     {1 1 2}      ; options for the number of channels  
node_arc_raw_format      {1 2 1}      ; options for the raw format STREAM_S16, STREAM_FP32  
end
```



# Graph (a text file : manual input or generated by a GUI)

## Nodes

```
arm_stream_filter 0
parameters
  1 u8; 0
  2 u8; 2 0
  5 h16; 1231 1D28 1231 63E8 D475
  5 h16; 1231 0B34 1231 2470 9821
end
```

node name instance index Boot preset,  
Options : Memory allocation, pre/post processing script,  
Dedicated architecture, processor (or any), priority, trace verbose level  
Memory mapping of each segment

## Arcs

```
arm_stream_filter 0 1
arm_stream_detector 0 0
```

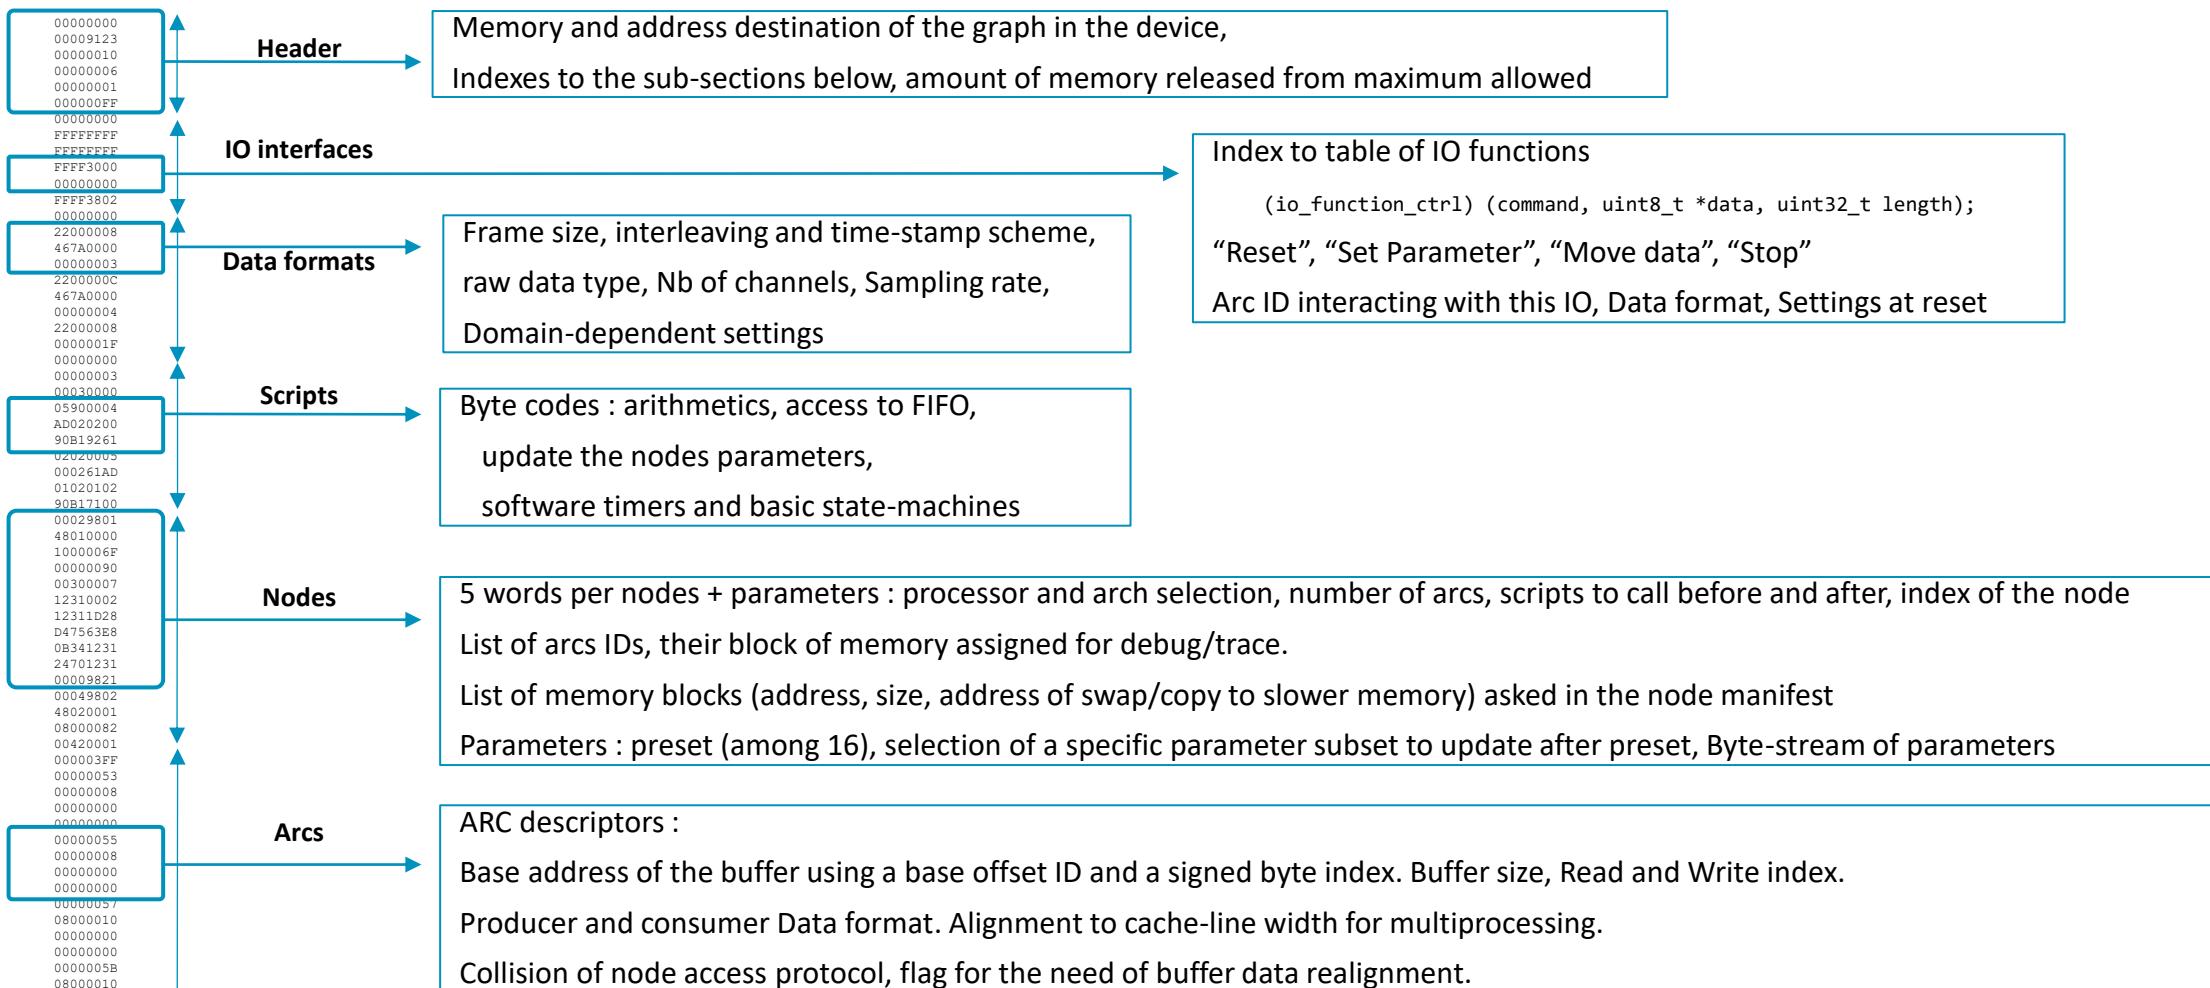
node name instance index arc output index  
Node name instance index arc input index

```
arc_input 0 0
arm_stream_filter 0 0 0
```

Arcs at the boundary of the graph  
node it is connected to



# “Compiled” Graph (used by the scheduler)



# Processor manifest : memory mapping

```
;=====
; TOP MANIFEST :
;   - paths to the files
;   - shared memories
;   - list of processors and their private memories and IOs
;   - list of the nodes installed in each processor and/or
;     architectures
;=====
6           six file paths
```

```
0 ../../stream_platform/
1 ../../stream_platform/computer/manifest/
2 ../../stream_nodes/arm/
3 ../../stream_nodes/signal-processingFR/
4 ../../stream_nodes/bitbank/
5 ../../stream_nodes/elm-lang/
```

```
;=====
3           number of processors
2           number of shared memory

;   processor and architecture ID are in the range [1..7]
;   Access      0 data/prog R/W, 1 data R, 2 R/W, 3 Prog, 4 R/W
;   Speed       0 slow/best effort, 1 fast/internal, 2 TCM
;   Type        0 Static, 1 Retention, 2 scratch mem

;---MEMID SIZE A S T   Comments
0    8000  1 0 0    shared memory
3    1000  1 0 1    simulates shared retention memory
```

```
;=====
; Processor #1 - architecture #1, two processors
; proc ID, arch ID, main proc, nb mem, service mask, I/O
1    1    1    2    15   7

;---MEMID SIZE A S T   Comments
1    1000  1 2 0    simulates DTCM
2    1000  1 2 1    simulates ITCM
```

```
;---IO AFFINITY WITH PROCESSOR 1---
,Path      Manifest    IO_AL_idx  Comments
1 io_platform_data_in_0.txt    0    application proc
1 io_platform_data_in_1.txt    1    application proc
1 io_platform_analog_sensor_0.txt 2    ADC
1 io_platform_audio_in_0.txt   4    microphone
1 io_platform_line_out_0.txt   6    audio out stereo
1 io_platform_gpio_out_0.txt   7    GPIO/LED
1 io_platform_data_out_0.txt   9    application proc
```

**Shared memory declaration**  
**Number of memory banks below, bit-field of  
 installed compute libraries**

**Private memory declaration**

**List of streams the processor can activate**

**List of nodes mapped to processors**

```
... (continued) ...
;=====
; Processor #2
; proc ID, arch ID, main proc, nb mem, service mask, I/O
2    1    0    2    15   2

;---MEMID SIZE A S T   Comments
1    1000  1 2 0    index 1/2 point to different physical addresses
2    1000  1 2 1    index 1/2 point to different physical addresses

;---IO AFFINITY WITH PROCESSOR 2---
;Path      IO Manifest    IO_AL_idx  Comments
1 io_platform_data_in_0.txt    0    shared with Proc 1
1 io_platform_motion_in_0.txt 3    accelero=gyro
```

```
;=====
; Processor #3 - new architecture, one processor
; proc ID, arch ID, main proc, nb mem, service mask, I/O
1    2    0    0    15   2

;---MEMID SIZE A S T   Comments
;---IO AFFINITY WITH PROCESSOR 2---
;Path      IO Manifest    IO_AL_idx  Comments
1 io_platform_2d_in_0.txt    5    camera
1 io_platform_gpio_out_1.txt 8    GPIO/PWM
```

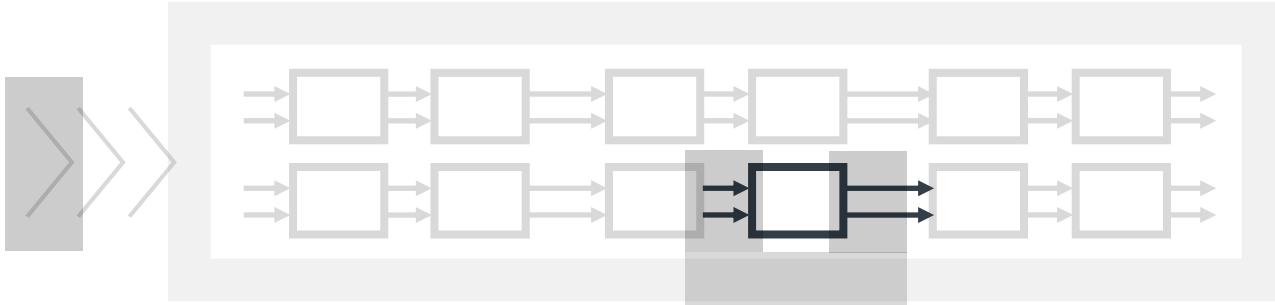
```
;===== ALL NODES =====
; scheduler algorithm :
; if the node archID > 0 then check compatibility with processor archID and exit
; if the node procID > 0 then check compatibility with processor procID and exit
;

; Path          Node Manifest          PROC ARCH | ID
2   script/node_manifest_script.txt  0 0 | 1 runs everywhere
2   router/node_manifest_router.txt 0 1 | 2 SMP on archID-1
2   amplifier/node_manifest_amplifier.txt 0 1 | 3
2   filter/node_manifest_filter.txt 0 1 | 4
2   modulator/node_manifest_modulator.txt 0 1 | 5
2   demodulator/node_manifest_demodulator.txt 0 1 | 6
3   detector/node_manifest_detector.txt 0 1 | 7
3   resampler/node_manifest_resampler.txt 0 1 | 8
3   compressor/node_manifest_compressor.txt 0 1 | 9
3   decompressor/node_manifest_decompressor.txt 0 1 | 10
4   JPEGENC/node_manifest_bitbank_JPEGENC.txt 0 1 | 11
5   TJpgDec/node_manifest_TjpgDec.txt 0 1 | 12
;
2   filter2D/node_manifest_filter2D.txt 2 2 | 13 only archID-2
3   detector2D/node_manifest_detector2D.txt 0 2 | 14 single processor
;
end ; the platform manifest ends here
=====
```



# Manifests of interfaces for Graph-I/Os

3



## 3 Graph-I/O interfaces :

a text file (readable syntax)

done once at platform manufacturing

```
io_platform_sensor_in_0          ; name for the tools
analog_in                         ; domain name, unit: dB, Vrms, mV/Gauss, dps, kWh, ...
                                    ; options for the raw arithmetics computation format here STREAM_S16
io_commander0_servant1 1          ; commander=0 servant=1 (default is servant)
io_buffer_allocation   2.0 1       ; default is 0, which means the buffer is declared outside of the graph, VID 1
io_direction_rx0tx1    1           ; direction of the stream 0:input 1:output from graph point of view
io_raw_format          {1 17}      ; multichannel interleaved (0), deinterleaved by frame-size (1) + options for the number of channels
io_nb_channels         {1 1 2}
io_frame_length        {1 2 16}
io_subtype_units       104          ; depending on the domain. Here Units_Vrms of the "general" domain (0 = any or undefined)
io_anologscale         0.55         ; 0.55V is corresponding to full-scale (0x7FFF or 1.0f) with the default setting
io_sampling_rate       {1 16000 44100 48000} ; sampling rate options (enumeration in Hz)
io_rescale_factor      12.24 -44.3  ; [1/a off] analog_input = invinterpa x ((samples/Full_Scale_Digital) - interpooff)
end
```



# Graph API (one entry-point to the scheduler)

## 1) Graph interpreter interface for the application :

```
void arm_graph_interpreter (uint32_t command, arm_stream_instance_t *S, uint8_t *data, uint32_t size)
```

Commands : reset the graph, execute, check boundary FIFO filling state and move data in/out, update the use-case

Instance : structure of pointers to the graph, to the installed nodes and application callbacks, to the data stream interfaces functions (below), control fields and static memory of the scheduler instance.

## 2) Stream interfaces used by the scheduler to initiate data moves (abstraction layer of the BSP):

```
void (io_function_ctrl) (uint32_t command, uint8_t *data, uint32_t length);
```

Commands : set buffer, set parameters, data move, stop

## 3) One callback, after data moves (to update the FIFO descriptors) :

```
void arm_graph_interpreter_io_ack (uint8_t fw_io_idx, uint8_t *data, uint32_t data_size)
```

## 4) One prototype for all nodes :

```
void node_XXXX (uint32_t command, void *instance, void *data, uint32_t *status)
```

command = reset, set parameters, run, stop

“manual” flow control from the application or  
let the scheduler call the BSP



Abstraction layer of IOs : data-move and settings + callback to set the FIFO  
or  
Data move from the application with same functions for FIFO setting



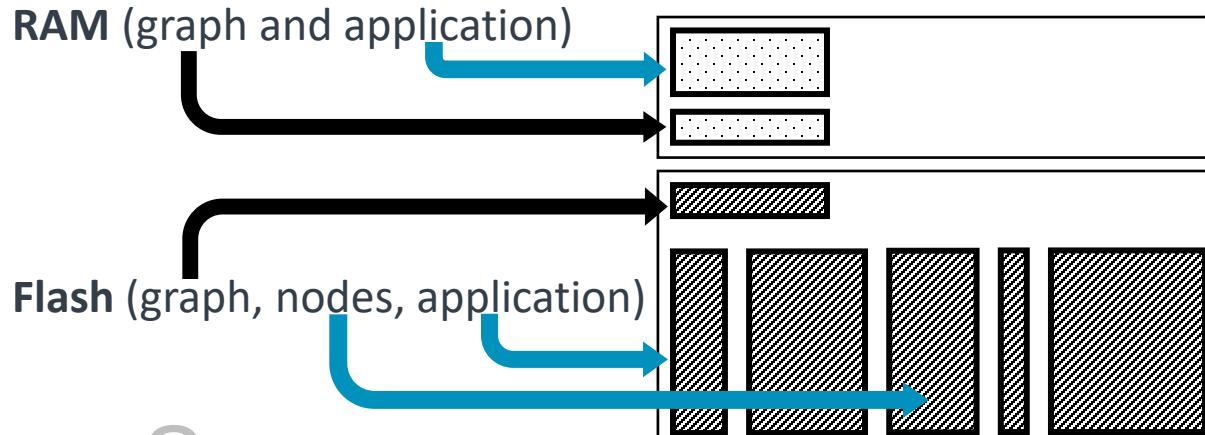
# Small memory footprint for LoRA

Remote sensors connected through LoRA have a data rate as low as 50Bytes/s

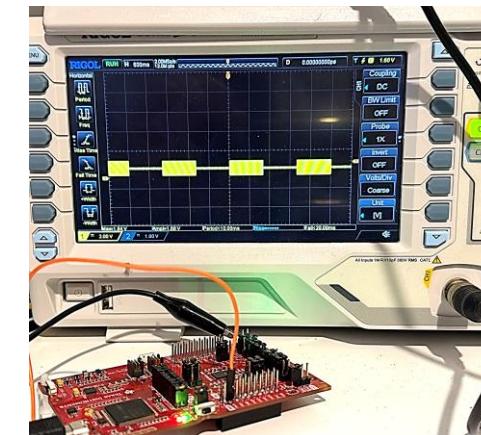
A graph size of two nodes (+ their respective parameters and a script) is in the 500Bytes range

**An interpreter eliminates the risk of malware injection during firmware updates**

A memory map of the LoRA device :



Filter and detector nodes with 1kB-RAM



# Graph with embedded scripts

A graph can incorporate nodes with interpreted code using basic integer/float arithmetics.

The instruction “CALLSYS” gives access to nodes (set/read parameters), arcs (read/write, check access time-stamps), application callbacks, etc..

The script interpreter is consuming less than 100 Bytes of stack memory.

**Why would you need Python for very simple operations similar to the ones of pocket calculator ?**

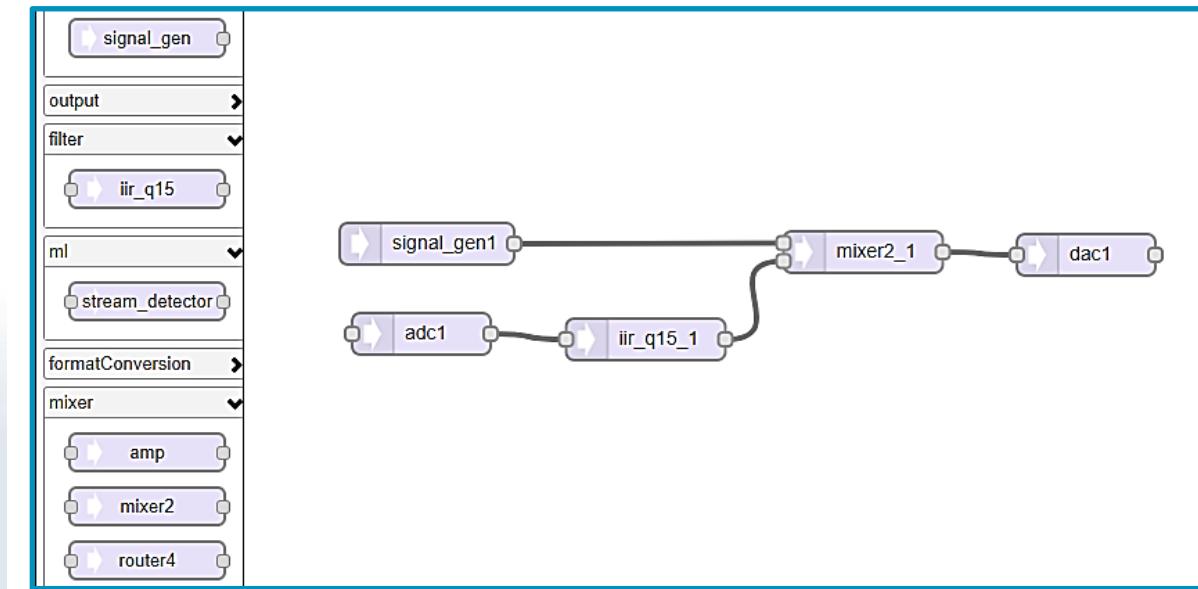
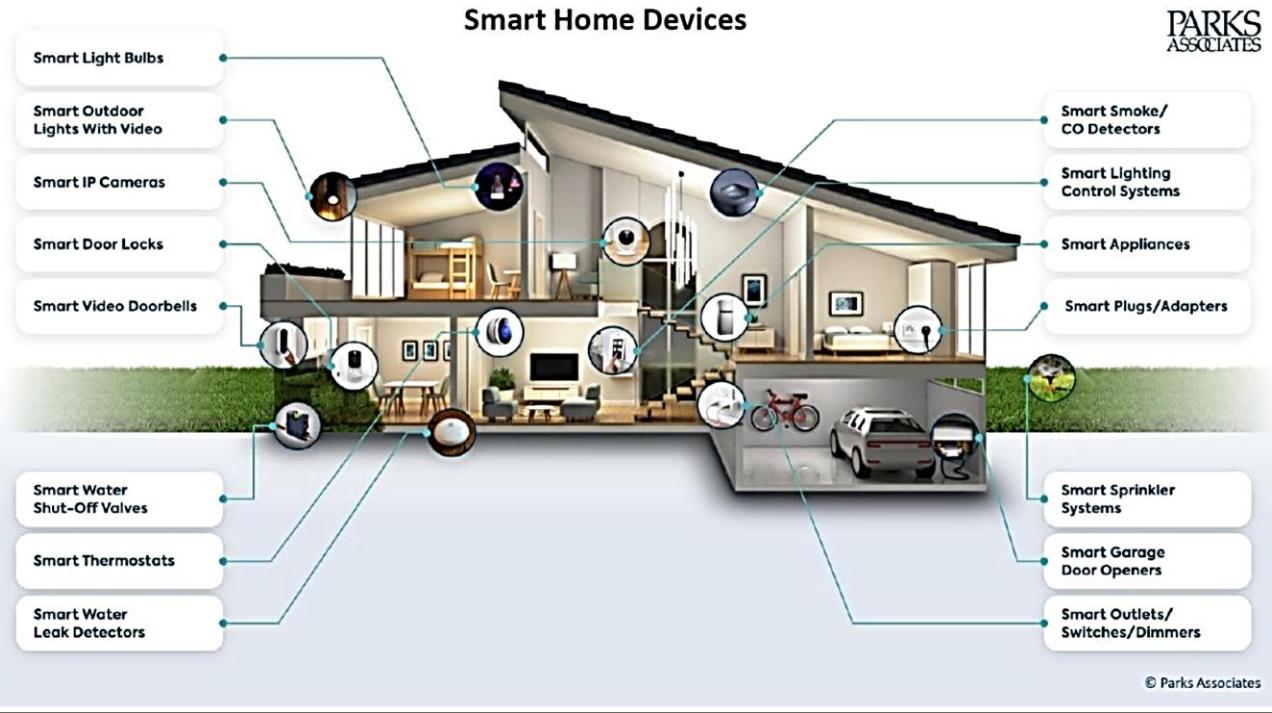
## Examples of instructions

r6 = r5 + 3.14	r6 = r5 + 3.14
S = r6	push the result on stack
test r6 == r5 - r4	test if r6 == ( r5 - r4 )
if_yes call label_xyz	conditional call
r3 8 15  = r2	bit-field load of r2 to the 2nd byte of r3
r3 = r4[12]	gather load r3 = r4[12]
banz L1 r2	decrement r2 and branch if not zero
Other operations: div,or,nor,and,xor,shr,shl,set,clr,max,min,amax,amin,norm,addmodulo	
syscall 4 r1 r5 r10	system call #4 with three parameters



# Next steps : low-code for smart-home sensors

Do we need a complex programming environment to drag and drop software components from a Store ?



GUI proof of concept using node-RED

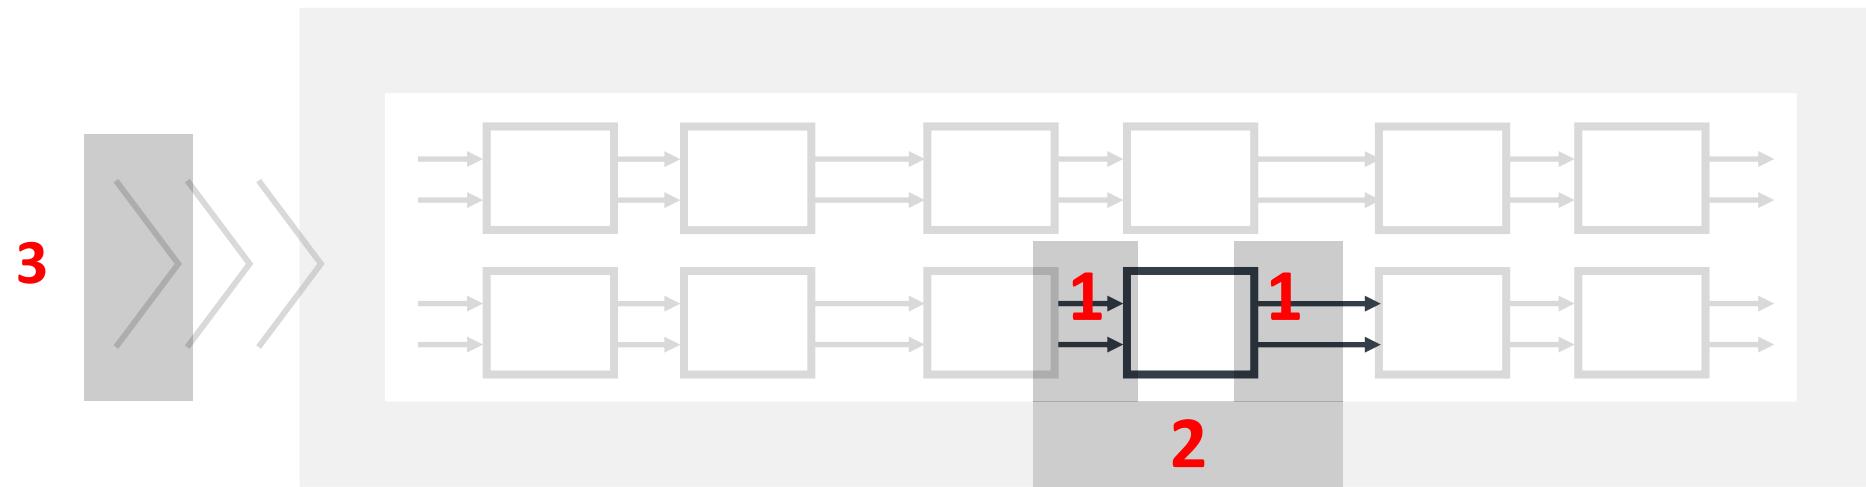


# Next steps (1Q26)

- Code validation on different platforms : Renesas RA8, Infineon E84, ALIF E7 (next: ST C011, WIO SAMD51, TI MSPM01306, NXP MCXC041).
- New nodes (currently just prototypes) : audio encoder/decoder, format converter (next: signal generator, image motion detector).
- Validation of the multiprocessing capability with low-power + high-performance cores executing the same graph
- Extension of the compute services : from current filters to FFT and math operators



# BACKUP



# arm

Thank You

Danke

Gracias

Grazie

謝謝

ありがとう

Asante

Merci

감사합니다

ধন্যবাদ

Kiitos

شکرًا

ধন্যবাদ

ଧାରା

ଧନ୍ୟବାଦମୁଲୁ





The Arm trademarks featured in this presentation are registered trademarks or trademarks of Arm Limited (or its subsidiaries) in the US and/or elsewhere. All rights reserved. All other marks featured may be trademarks of their respective owners.

[www.arm.com/company/policies/trademarks](http://www.arm.com/company/policies/trademarks)

